

PIV Project

Group 9	
Alexandre Cachera	116285
Tiago Nunes Couto	113311
Bernardo Silva Costa	102777
Gabriel Badan	116537

1 Motivations

This project aims to estimate geometric transformations between a reference image and a sequence of images using classical and modern computer vision techniques. The project is structured in two parts: (1) *Part 1*: 2D homography estimation from grayscale images, and (2) *Part 2*: rigid 3D transformation (rotation and translation) from RGB-D data. These techniques find applications in object tracking, 3D reconstruction, robotic localization, and augmented reality.

2 Part 1: Planar Object Tracking via Homography

This section describes the implementation of a 2D planar tracking system that estimates the Homography (H) between a reference template and a sequence of video frames.

2.1 Algorithmic Details

2.1.1 Feature Extraction and Matching

The system utilizes the **Scale-Invariant Feature Transform (SIFT)** [1] to detect interest points robust to scale

and rotation, and extract 128-dimensional invariant descriptors. We selected SIFT over simpler alternatives (Harris corners, ORB) because the document occupies varying image regions as the camera moves, requiring scale-invariant detection. Additionally, SIFT handles arbitrary document orientations through its rotation-invariant formulation, and unlike learning-based methods, it requires no training data and remains fully interpretable. *SIFT Parameters*: OpenCV defaults ($\sigma = 1.6$, standard octaves and scales).

Then, feature matching is performed between the template and each frame using the following strategy:

We use Euclidean distance metric combined with **Lowe's ratio test** to filter ambiguous matches. For each descriptor in the frame, we find the 2 nearest neighbors in the template. If $\text{distance}(\text{NN}_1) < 0.75 \times \text{distance}(\text{NN}_2)$, we accept the match.

The selection of the 0.75 threshold for the distance ratio test is grounded in D. Lowe's original empirical analysis [1]. This test aims to eliminate ambiguous matches by comparing the Euclidean distance of the nearest neighbor to that of the second-nearest neighbor. A ratio exceeding 0.75 indicates that the two best candidates are similarly close, suggesting that the feature lacks distinctiveness—often due to repetitive patterns or background clutter.

We chose this specific value because it offers an optimal trade-off between precision and recall. **Higher thresholds** (e.g., 0.9) tend to retain too many incorrect matches where the descriptor is not unique enough, while **stricter thresholds** (e.g., 0.5) are too aggressive and discard valid correspondences, particularly in regions with lower texture variations. Thus, 0.75 effectively filters out unreliable keypoints while maintaining a dense set of features for the subsequent RANSAC estimation.

2.1.2 Homography Estimation (DLT with Normalization)

The relationship between planar points in the template ($\mathbf{x}$) and the frame ($\mathbf{x}'$) is modeled by a 3×3 Homography matrix H such that $\mathbf{x}' \sim H\mathbf{x}$. The **Direct Linear Transform (DLT)** algorithm is used to solve for H using Singular Value Decomposition (SVD).

Data Normalization A crucial step implemented is isotropic normalization [2]. Before applying DLT, points are translated so their centroid is at the origin and scaled so the average distance from the origin is $\sqrt{2}$. This improves the numerical stability of the SVD.

$$T_{norm} = \begin{bmatrix} s & 0 & -s \cdot c_x \\ 0 & s & -s \cdot c_y \\ 0 & 0 & 1 \end{bmatrix} \quad (1)$$

Matrix A Construction For N correspondence points $(x_i, y_i) \rightarrow (u_i, v_i)$, create a $2N \times 9$ matrix:

$$A_i = \begin{pmatrix} -x_i & -y_i & -1 & 0 & 0 & 0 & x_i u_i & y_i u_i & u_i \\ 0 & 0 & 0 & -x_i & -y_i & -1 & x_i v_i & y_i v_i & v_i \end{pmatrix}$$

SVD and H Extraction The solution is found via SVD: $A = U\Sigma V^T$, where H is reshaped from the last column of V (corresponding to the smallest singular value).

Denormalization The final Homography is denormalized via:

$$H = T_{dst}^{-1} H_{norm} T_{src} \quad (2)$$

2.1.3 Robust Estimation (RANSAC)

Even after Lowe's ratio test, some incorrect correspondences persist—particularly in low-texture regions or near occlusion boundaries. A direct least-squares solution to the homography would be severely biased by these outliers, potentially producing unusable results. RANSAC addresses this by iteratively searching for the largest subset of geometrically consistent matches.

Our implementation uses the following parameters: **Iterations** (5000), **Sample Size** (4 points, the algebraic minimum for homography), **Inlier Threshold** (3.0 pixels), and **Refinement** (after identifying the best consensus set, H is re-computed using all inliers for improved accuracy).

Inlier Threshold Selection The reprojection error threshold was set to 3.0 pixels based on an empirical analysis of the localization error inherent in our video data. At the typical working distances and camera resolutions used in this project, a deviation of 3.0 pixels corresponds to approximately 1–2 mm on the physical document surface. This specific value allows the system to accommodate the sensor noise floor and minor feature localization inaccuracies, while effectively rejecting correspondence errors caused by occlusions or texture ambiguities. A lower threshold would discard valid matches due to discretization noise, whereas a higher threshold would admit outliers, degrading the stability of the estimated homography.

Iteration Strategy To ensure robust model estimation without compromising runtime performance, we implemented an iteration scheme with 5000 iterations. The theoretical number of iterations needed to achieve a confidence probability p is:

$$k = \frac{\log(1-p)}{\log(1-w^s)} \quad (3)$$

where w is the inlier ratio and $s = 4$ is the sample size. With $p = 0.99$, this ensures a 99% chance of sampling a minimal set of outlier-free points. The 5000-iteration budget provides a computational margin for difficult frames where the inlier ratio may be low (e.g., during severe motion blur or occlusion).

2.2 Experimental Setup and Evaluation

Dataset We used the TAAG dataset, which provides a reference template image and a sequence of frames with varying perspectives and occlusions.

Evaluation metrics We considered the *number of inliers* as an indicator of matching quality and *robustness* as an ability to handle partial occlusions and significant transformations.

Execution Pipeline The implementation is organized in two stages: (1) **Template Initialization:** SIFT extracts features from the template image and saves them to `template_features.mat`. (2) **Frame-by-Frame Processing:** For each frame, SIFT features are extracted and stored; feature matching via Lowe’s ratio test identifies correspondences; RANSAC (5000 iterations, 3.0 pixel threshold) estimates the best homography; the final H is refined using all inliers and saved as `homography_NNNN.mat`. Frames with fewer than 4 matches are skipped.

2.3 Example Results

We evaluated the algorithm on a 14-frame sequence with varying conditions.




Figure 1: Successful tracking under perspective change.




Figure 2: Robust tracking despite partial occlusion—RANSAC rejects outliers.



Figure 3: Failure case: insufficient correspondences produce incorrect homography.

Quantitative Analysis Table 1 summarizes the results across the 14 processed frames, categorized by visibility conditions.

Table 1: Part 1: Homography estimation statistics by frame type.

Type	Frames	Matches	Inliers (%)	Error (px)
Clear	5–8	1679–2235	58–72%	1.1–1.5
Partial occl.	2–4, 9–12	434–2341	21–57%	1.5–1.6
Failure	1, 13–14	64–228	10–23%	0.5–142*
Mean	–	1158	41%	1.4**

*Frame 0001 exhibits extremely high error (142 px), indicating a degenerate homography due to poor feature distribution.

**Mean error computed excluding frame 0001.

Key Observations

- **Clear visibility** (frames 5–8): High match counts (1600–2200) with inlier ratios above 58%, producing stable homographies with sub-pixel accuracy.
- **Partial occlusions** (frames 2–4, 9–12): RANSAC successfully filters occluded features, maintaining reasonable accuracy despite reduced inlier ratios.
- **Failure cases** (frames 1, 13–14): Low match counts (<230) lead to unreliable estimation. Frame 0001 pro-

duces a degenerate result due to poor initial feature distribution.

3 Part 2: 3D Rigid Body Motion Estimation

This section details the estimation of the camera’s 3D trajectory (Odometry) by calculating the Rigid Body Transformation (R, T) between the reference template and each frame.

3.1 Algorithmic Details

3.1.1 Feature Extraction and Matching

Feature matching uses SIFT on RGB images to identify 2D-2D correspondences, applies Lowe’s ratio test (threshold = 0.75), and recovers descriptors robust to lighting changes.

3.1.2 3D Point Generation

Unlike Part 1, this module utilizes RGB-D data. 2D key-points (u, v) are back-projected into 3D space (X, Y, Z) using the corresponding depth map and the Pinhole Camera Model intrinsics matrix K :

$$Z = \text{depth}(v, u) \quad (4)$$

$$X = \frac{(u - c_x) \cdot Z}{f_x}, \quad Y = \frac{(v - c_y) \cdot Z}{f_y} \quad (5)$$

Points with invalid depth ($Z \leq 0$ or NaN) are discarded.

3.1.3 Orthogonal Procrustes Analysis

To find the rotation R and translation T that align two clouds of 3D points ($P_B \approx R \cdot P_A + T$), we use the Orthogonal Procrustes method [3]. Unlike iterative approaches such as ICP (Iterative Closest Point), Procrustes assumes that point correspondences are already established (via SIFT matching). This allows a direct closed-form solution rather than an iterative search, resulting in faster computation and guaranteed convergence. The method proceeds as follows:

1. Compute centroids: $\bar{P}_A = \frac{1}{N} \sum P_A^{(i)}, \bar{P}_B = \frac{1}{N} \sum P_B^{(i)}$

2. Center both point clouds: $P'_A = P_A - \bar{P}_A, P'_B = P_B - \bar{P}_B$

3. Compute cross-covariance matrix: $H = (P'_A)^T \cdot P'_B$

4. Decompose via SVD: $H = U \Sigma V^T$

5. Compute rotation: $R = V U^T$

6. If $\det(R) < 0$, negate the last column of V and recompute R [4]

7. Compute translation: $T = \bar{P}_B - R \cdot \bar{P}_A$

3.1.4 RANSAC for 3D Alignment

Depth sensors introduce noise near object boundaries, and SIFT matches can fail in textureless or specular regions. Applying Procrustes directly to all correspondences would be corrupted by these outliers. We therefore wrap the Procrustes solver in a RANSAC loop with: **Iterations** (2000), **Sample Size** (4 points, providing redundancy over the theoretical minimum of 3), and **Threshold** (0.05 meters = 5 cm), applied to the Euclidean distance between transformed source and target points.

Threshold Selection The 5 cm threshold was determined through empirical analysis of our RGB-D sensor. Correctly matched points typically exhibit errors within this range due to sensor quantization, while mismatches or boundary artifacts produce errors exceeding 10 cm. After finding the largest consensus set, the final (R, T) is recomputed using all inliers.

3.2 Experimental Setup and Evaluation

Dataset The dataset comprises 1 RGB template with depth (templatedepth.mat), multiple RGB-D frames (frame_NNNN.mat), typical resolution of 640×480 pixels, and depth format as $H \times W$ matrix with values in millimeters.

Evaluation metrics Metrics include: **Number of Valid 3D Correspondences** (indicator of matching quality and depth density), **RANSAC Inliers** (ratio of robustness in pose estimation), and **Average 3D Error** (for inlier points).

Execution Pipeline The system iterates through pairs of RGB and Depth files: (1) Load RGB template with depth; (2) For each frame—extract and match SIFT features, back-project matched 2D points to 3D using depth maps, run RANSAC + Procrustes to estimate R, T , and store results in `transform_NNNN.mat` files.

3.3 Example Results

We evaluated the 3D rigid transformation estimation on a 12-frame RGB-D sequence. Figure 4 illustrates the complete pipeline output, showing the four processing stages: Input, Warped, Mask, and Cleaned result.



Figure 4: RGB-D pipeline visualization. From left to right: (1) Input frame, (2) Warped image after applying R, T , (3) Depth-based segmentation mask, (4) Cleaned document with occlusions removed.



Figure 5: Handling occlusions: the depth-based mask correctly identifies and removes the hand, preserving only the document plane.




Figure 6: Additional result demonstrating consistent alignment across frames.

Visual Analysis The pipeline successfully performs three key operations:

1. **3D Alignment:** The warped image shows the current frame transformed to align with the template coordinate system using the estimated (R, T) .
2. **Plane Segmentation:** The depth-based mask distinguishes the document plane from foreground occlusions (hands) and background objects by comparing measured depth against the expected document depth.
3. **Occlusion Removal:** The cleaned output preserves only pixels belonging to the document surface, effectively removing hands and other obstacles.

Quantitative Observations Across the 12 processed frames:

- All computed rotation matrices satisfy $R^T R = I$ and $\det(R) = 1$
- The depth threshold of 5 cm effectively separates document pixels from hand occlusions
- RANSAC convergence typically occurs within 500–1000 iterations (of the 2000 budget)

Failure Cases The system struggles when:

- Depth sensor returns invalid values (holes near object edges)
- The document surface is highly specular, causing depth dropout
- Fewer than 4 valid 3D correspondences remain after depth filtering

Result Quality The visual results show imperfect alignment in some frames. This is primarily due to weak feature matching between the template and the frames—the template exhibited limited overlap with most camera viewpoints, reducing the number of reliable correspondences available for pose estimation.

4 Discussion and Limitations

4.1 Comparison of Methods

2D Homography (Part 1) is efficient for tracking planar objects but fails on non-planar targets or when camera motion induces significant parallax that a planar model cannot represent. DLT normalization is critical for maintaining accuracy with large coordinates. **3D Procrustes (Part 2)** provides full 6-DoF pose estimation and requires depth data, making it dependent on depth sensor quality. The implementation correctly handles coordinate system reflections, a common pitfall in SVD-based estimation.

4.2 Identified Limitations

Part 1 Limitations include: limited to planar scenes, sensitivity to lighting changes (though SIFT mitigates this), and requirement for sufficient extractable features.

Part 2 Limitations include: critical dependency on depth data quality (noise, holes), degraded performance if valid 3D correspondences < 4 (RANSAC sample size), and sensitivity to calibration of K .

4.3 Conclusion

The implemented pipeline demonstrates a standard computer vision approach: Feature Extraction $\rightarrow$ Matching $\rightarrow$ Robust Model Estimation. The use of SVD for both DLT and Procrustes ensures mathematically rigorous solutions to the minimization problems.

References

- [1] D. G. Lowe, “Distinctive image features from scale-invariant keypoints,” *International Journal of Computer Vision*, vol. 60, no. 2, pp. 91–110, 2004.
- [2] R. Hartley and A. Zisserman, *Multiple View Geometry in Computer Vision*, 2nd ed. Cambridge University Press, 2003.
- [3] K. S. Arun, T. S. Huang, and S. D. Blostein, “Least-squares fitting of two 3-D point sets,” *IEEE Trans. Pattern Analysis and Machine Intelligence*, vol. PAMI-9, no. 5, pp. 698–700, 1987.
- [4] S. Umeyama, “Least-squares estimation of transformation parameters between two point patterns,” *IEEE Trans. Pattern Analysis and Machine Intelligence*, vol. 13, no. 4, pp. 376–380, 1991.